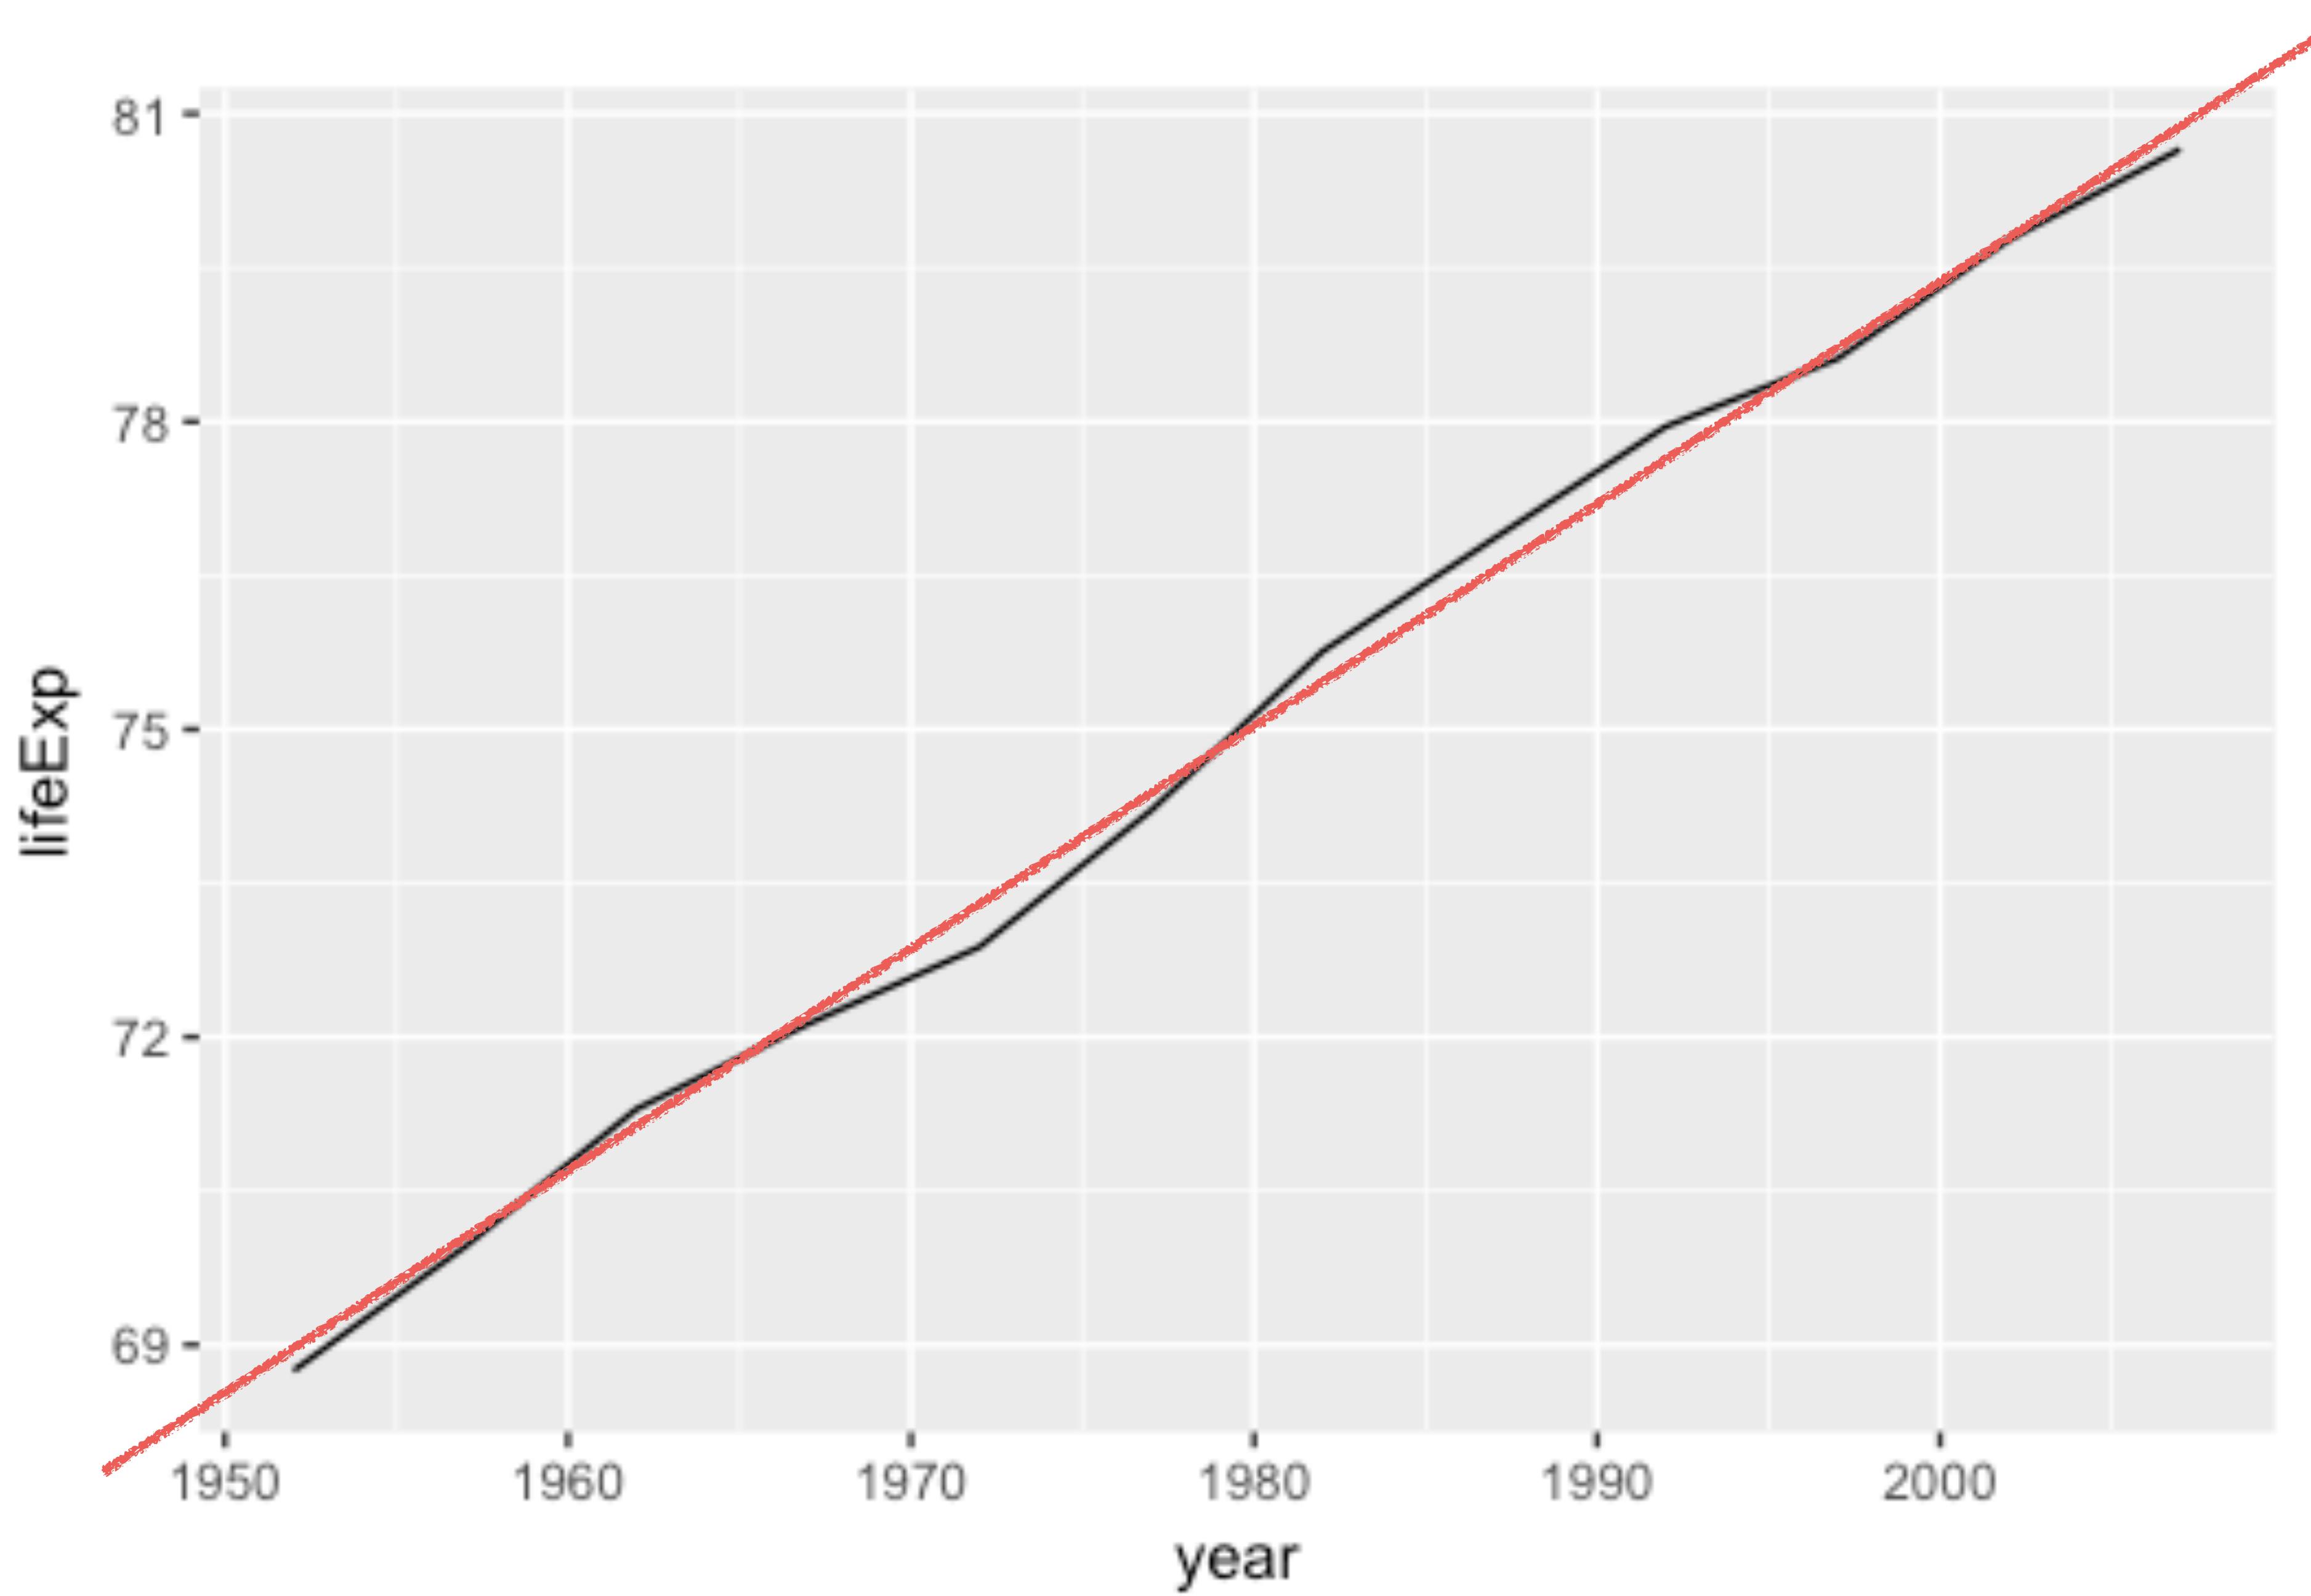


Modelling: the basics

Your Turn 1

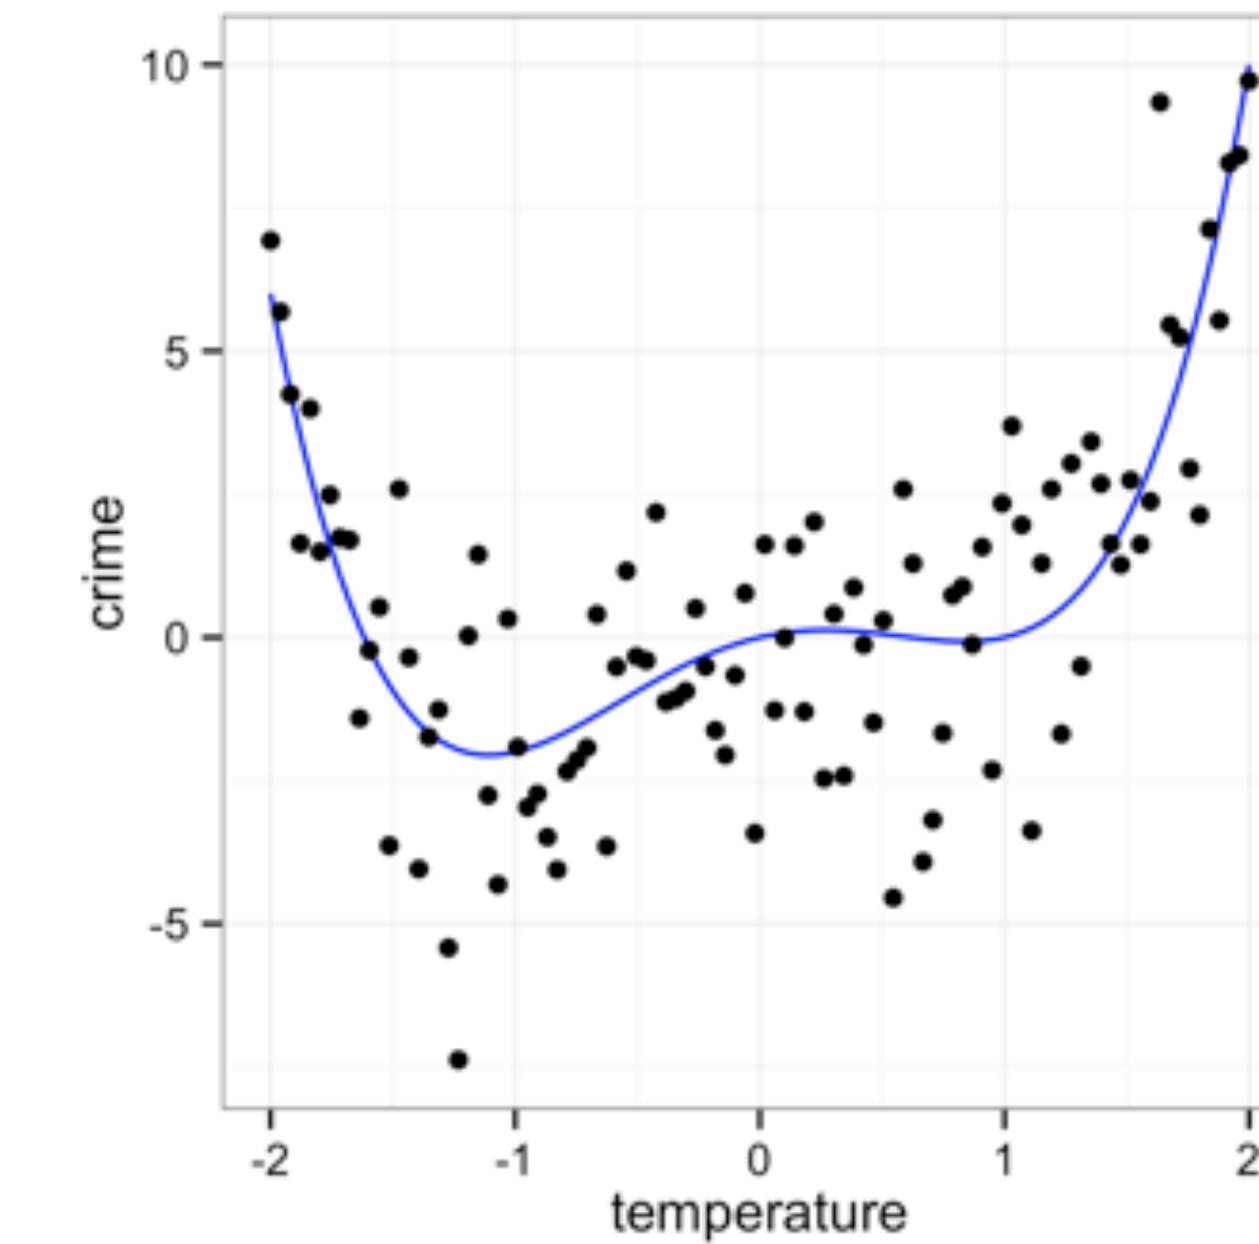
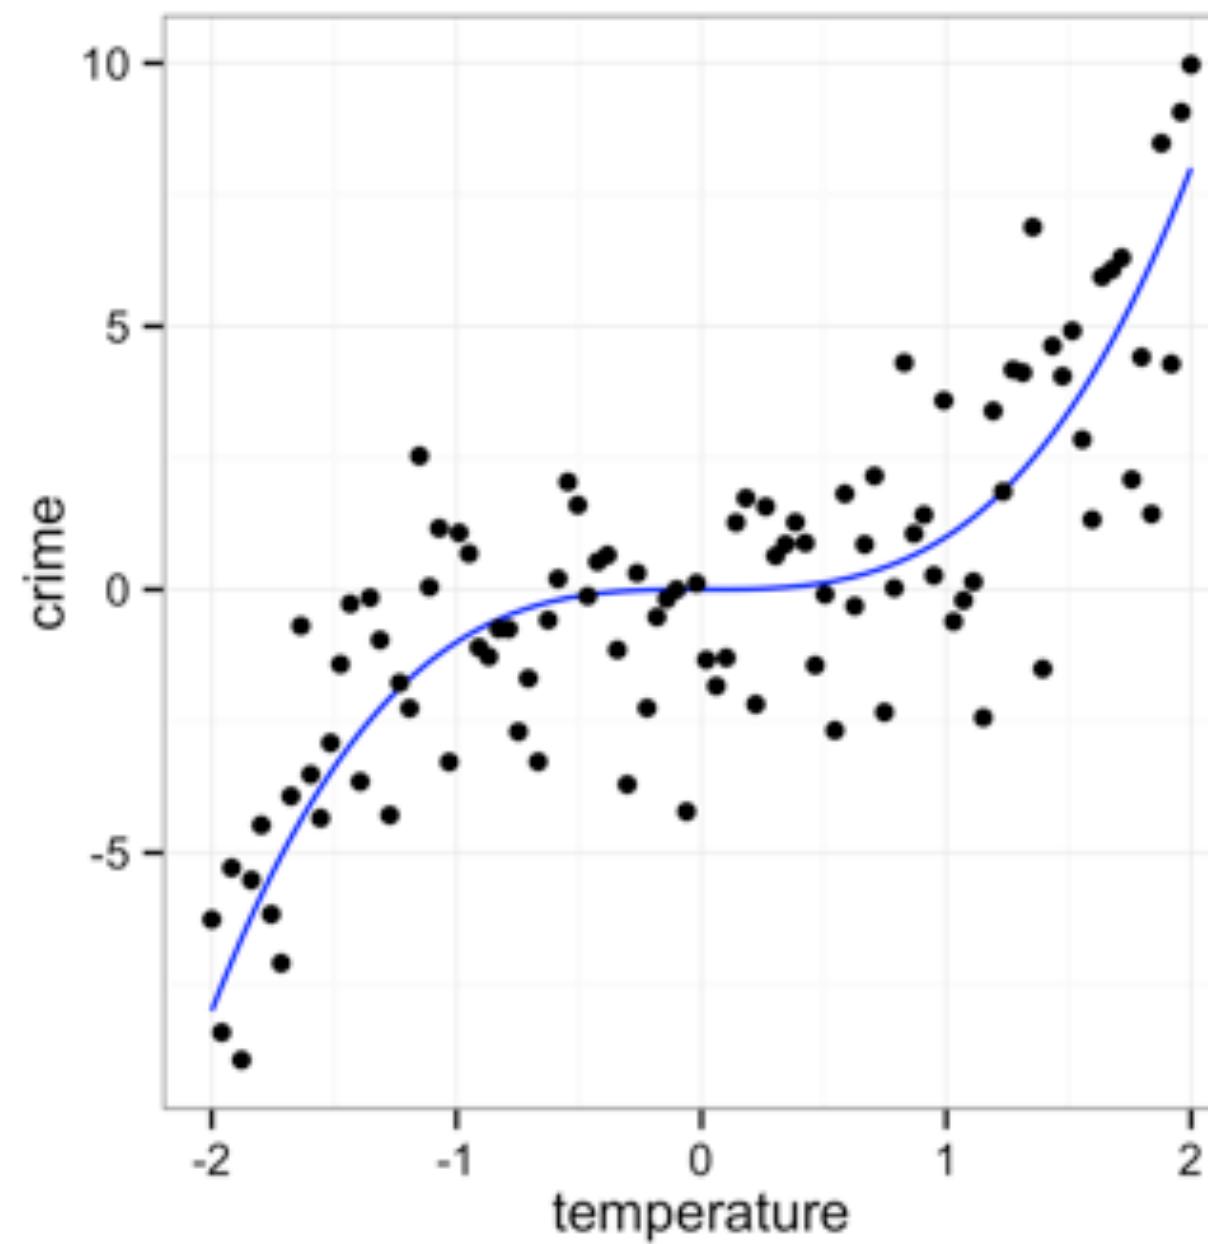
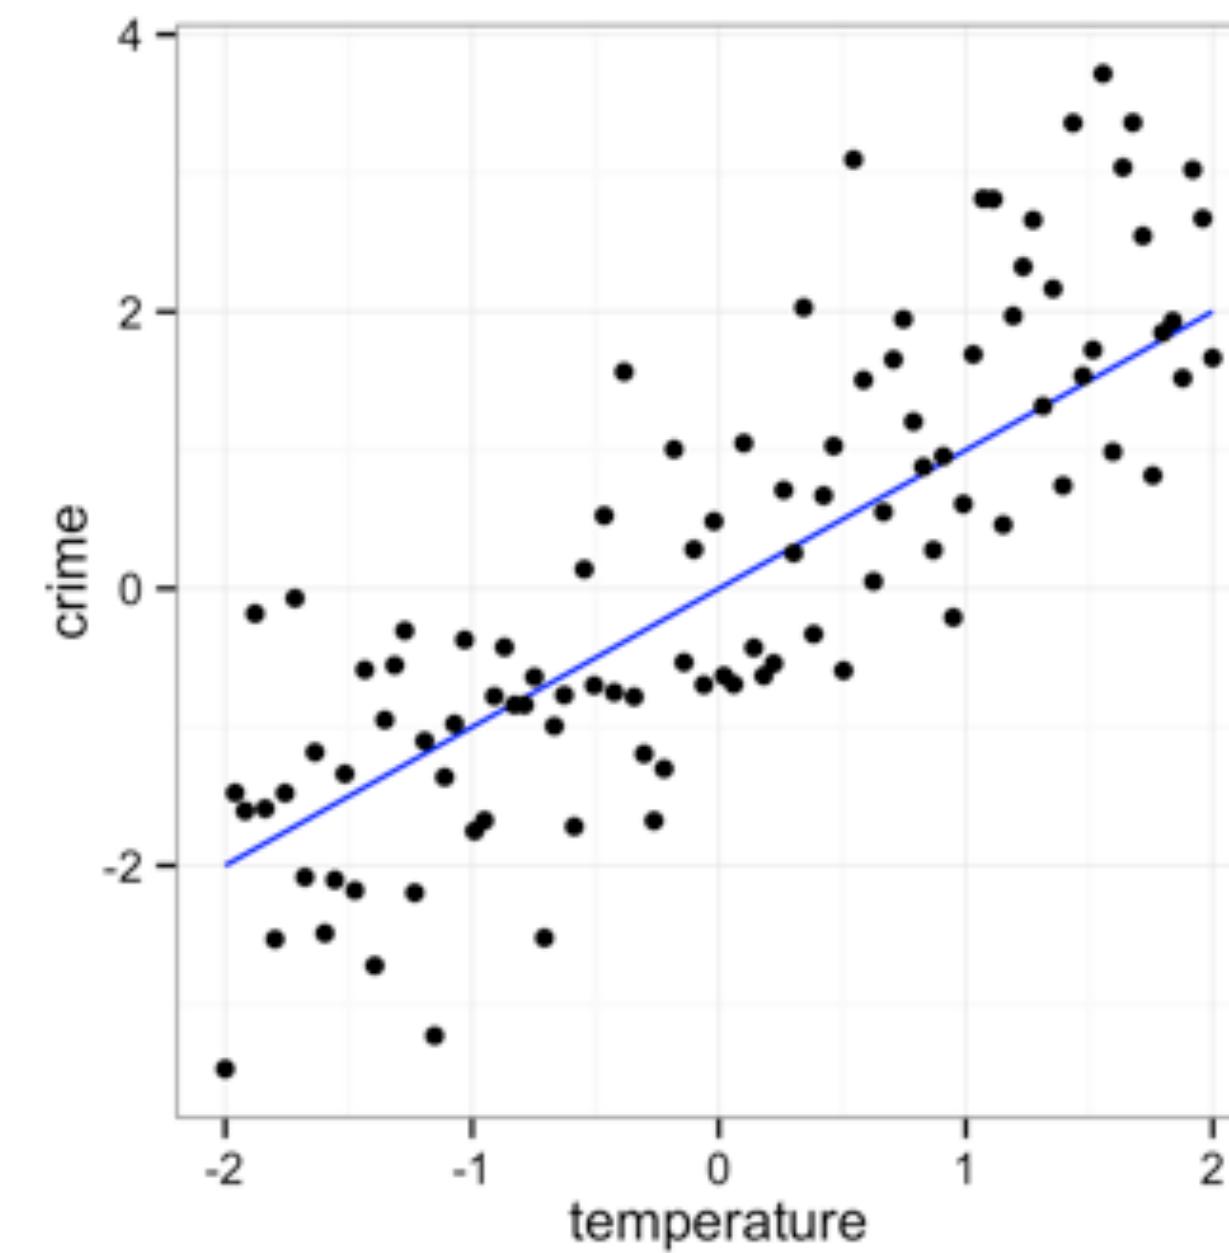
Take another look at life expectancy over time
for Canada.

What do you notice?



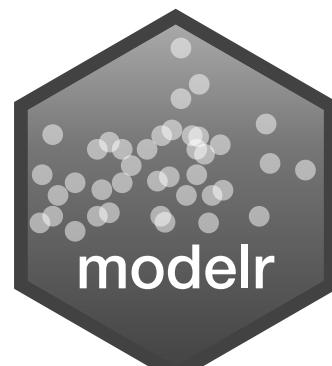
Models

A low dimensional description of a higher dimensional data set.



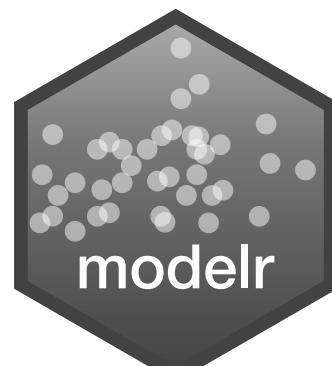
(Popular) modeling functions in R

function	package	fits
lm()	stats	linear models
glm()	stats	generalized linear models
gam()	mgcv	generalized additive models
glmnet()	glmnet	penalized linear models
rlm()	MASS	robust linear models
rpart()	rpart	trees
randomForest()	randomForest	random forests
xgboost()	xgboost	gradient boosting machines



(Popular) modeling functions in R

function	package	fits
lm()	stats	linear models
glm()	stats	generalized linear models
gam()	mgcv	generalized additive models
glmnet()	glmnet	penalized linear models
rlm()	MASS	robust linear models
rpart()	rpart	trees
randomForest()	randomForest	random forests
xgboost()	xgboost	gradient boosting machines



Im()

lm()

Fit a linear model to data

```
lm(lifeExp ~ year, data = canada)
```

A formula that describes
the model equation

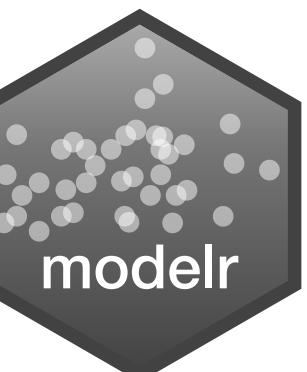
The data set

formulas

Formula only needs to include the response and predictors

$$y = \alpha + \beta x + \epsilon$$

$$\mathbf{y} \sim \mathbf{x}$$



Your Turn 2

Run the chunk to fit the model.

What kind of output do you get?

Explore `canada_lm` by clicking on it in the environment, what's inside the object?



```
Call:  
lm(formula = lifeExp ~ year, data = canada)
```

Coefficients:

(Intercept)	year
-358.3489	0.2189

Asking for the object prints a minimal summary

The screenshot shows the RStudio interface with the 'Environment' tab selected. The 'canada_lm' object is expanded, displaying its components and their values. The table has columns for Name, Type, and Value.

Name	Type	Value
canada_lm	list [12] (S3: lm)	List of length 12
coefficients	double [2]	-358.349 0.219
residuals	double [12]	-0.1338 -0.0182 0.2275 -0.0369 -0.38...
effects	double [12]	-2.59e+02 1.31e+01 2.56e-01 -7.88e-...
rank	integer [1]	2
fitted.values	double [12]	68.9 70.0 71.1 72.2 73.3 74.4 ...
assign	integer [2]	0 1
qr	list [5] (S3: qr)	List of length 5
df.residual	integer [1]	10
xlevels	list [0]	List of length 0
call	language	lm(formula = lifeExp ~ year, data = cana...
terms	formula	lifeExp ~ year
model	list [12 x 2] (S3: data.frame)	A data.frame with 12 rows and 2 columns

But the object has everything inside you might need to calculate useful summaries

summary(canada_lm)

Call:

```
lm(formula = lifeExp ~ year, data = canada)
```

Residuals:

Min	1Q	Median	3Q	Max
-0.3812	-0.1368	-0.0471	0.2481	0.3157

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-3.583e+02	8.252e+00	-43.42	1.01e-12 ***
year	2.189e-01	4.169e-03	52.50	1.52e-13 ***

Signif. codes: 0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 0.2492 on 10 degrees of freedom

Multiple R-squared: 0.9964, Adjusted R-squared: 0.996

F-statistic: 2757 on 1 and 10 DF, p-value: 1.521e-13

A more comprehensive summary, but not easy to work with.

broom

broom



Turns model output into data frames

```
# install.packages("tidyverse")
library(broom)
```



broom

Broom includes three functions which work for most types of models (and can be extended to more):

1. `tidy()` - returns model coefficients, stats
2. `glance()` - returns model diagnostics
3. `augment()` - returns predictions, residuals, and other raw values



tidy()

Returns useful model output as a data frame

```
canada_lm %>% tidy()
```

term	estimate	std.error	statistic	p.value
(Intercept)	-358.3488923	8.252132349	-43.42501	1.007334e-12
year	0.2188692	0.004168638	52.50378	1.520503e-13

2 rows

Easy to work with

Since output of tidy is a data frame we can use dplyr functions to manipulate it:

```
canada_lm %>% tidy() %>%  
  filter(term == "year") %>%  
  pull(estimate)
```

[1] 0.2188692

On average, life expectancy is increasing by 0.22 years per y

glance()

Returns common **model diagnostics** as a data frame

```
canada_lm %>% glance()
```

r.squared <dbl>	adj.r.squared <dbl>	sigma <dbl>	statistic <dbl>	p.value <dbl>	df <int>
0.9963855	0.9960241	0.2492483	2756.647	1.520503e-13	2

1 row | 1-6 of 11 columns



Your Turn 3

Look at the output of `canada_lm %>% glance()`

Use dplyr tools to extract the R-squared for the model.



```
canada_lm %>% glance() %>% pull(r.squared)
```

```
[1] 0.9963855
```

augment()

Returns data frame of model output related to original data points

```
canada_lm %>% augment()
```

lifeExp	year	.fitted	.se.fit	.resid	.hat
<dbl>	<int>	<dbl>	<dbl>	<dbl>	<dbl>
68.750	1952	68.88385	0.13534705	-0.13384615	0.29487179
69.960	1953	69.07919	0.11821353	-0.01810211	0.22494172
71.170	1954	70.27453	0.11821353	-0.10266464	0.16899767
72.150	1955	71.47986	0.11821353	-0.30933860	0.12703963
72.880	1972	73.26123	0.07845091	-0.38123077	0.09906760
74.210	1977	74.35558	0.07270261	-0.14557692	0.08508159
75.760	1982	75.44992	0.07270261	0.31007692	0.08508159
76.860	1987	76.54427	0.07845091	0.31573077	0.08906760

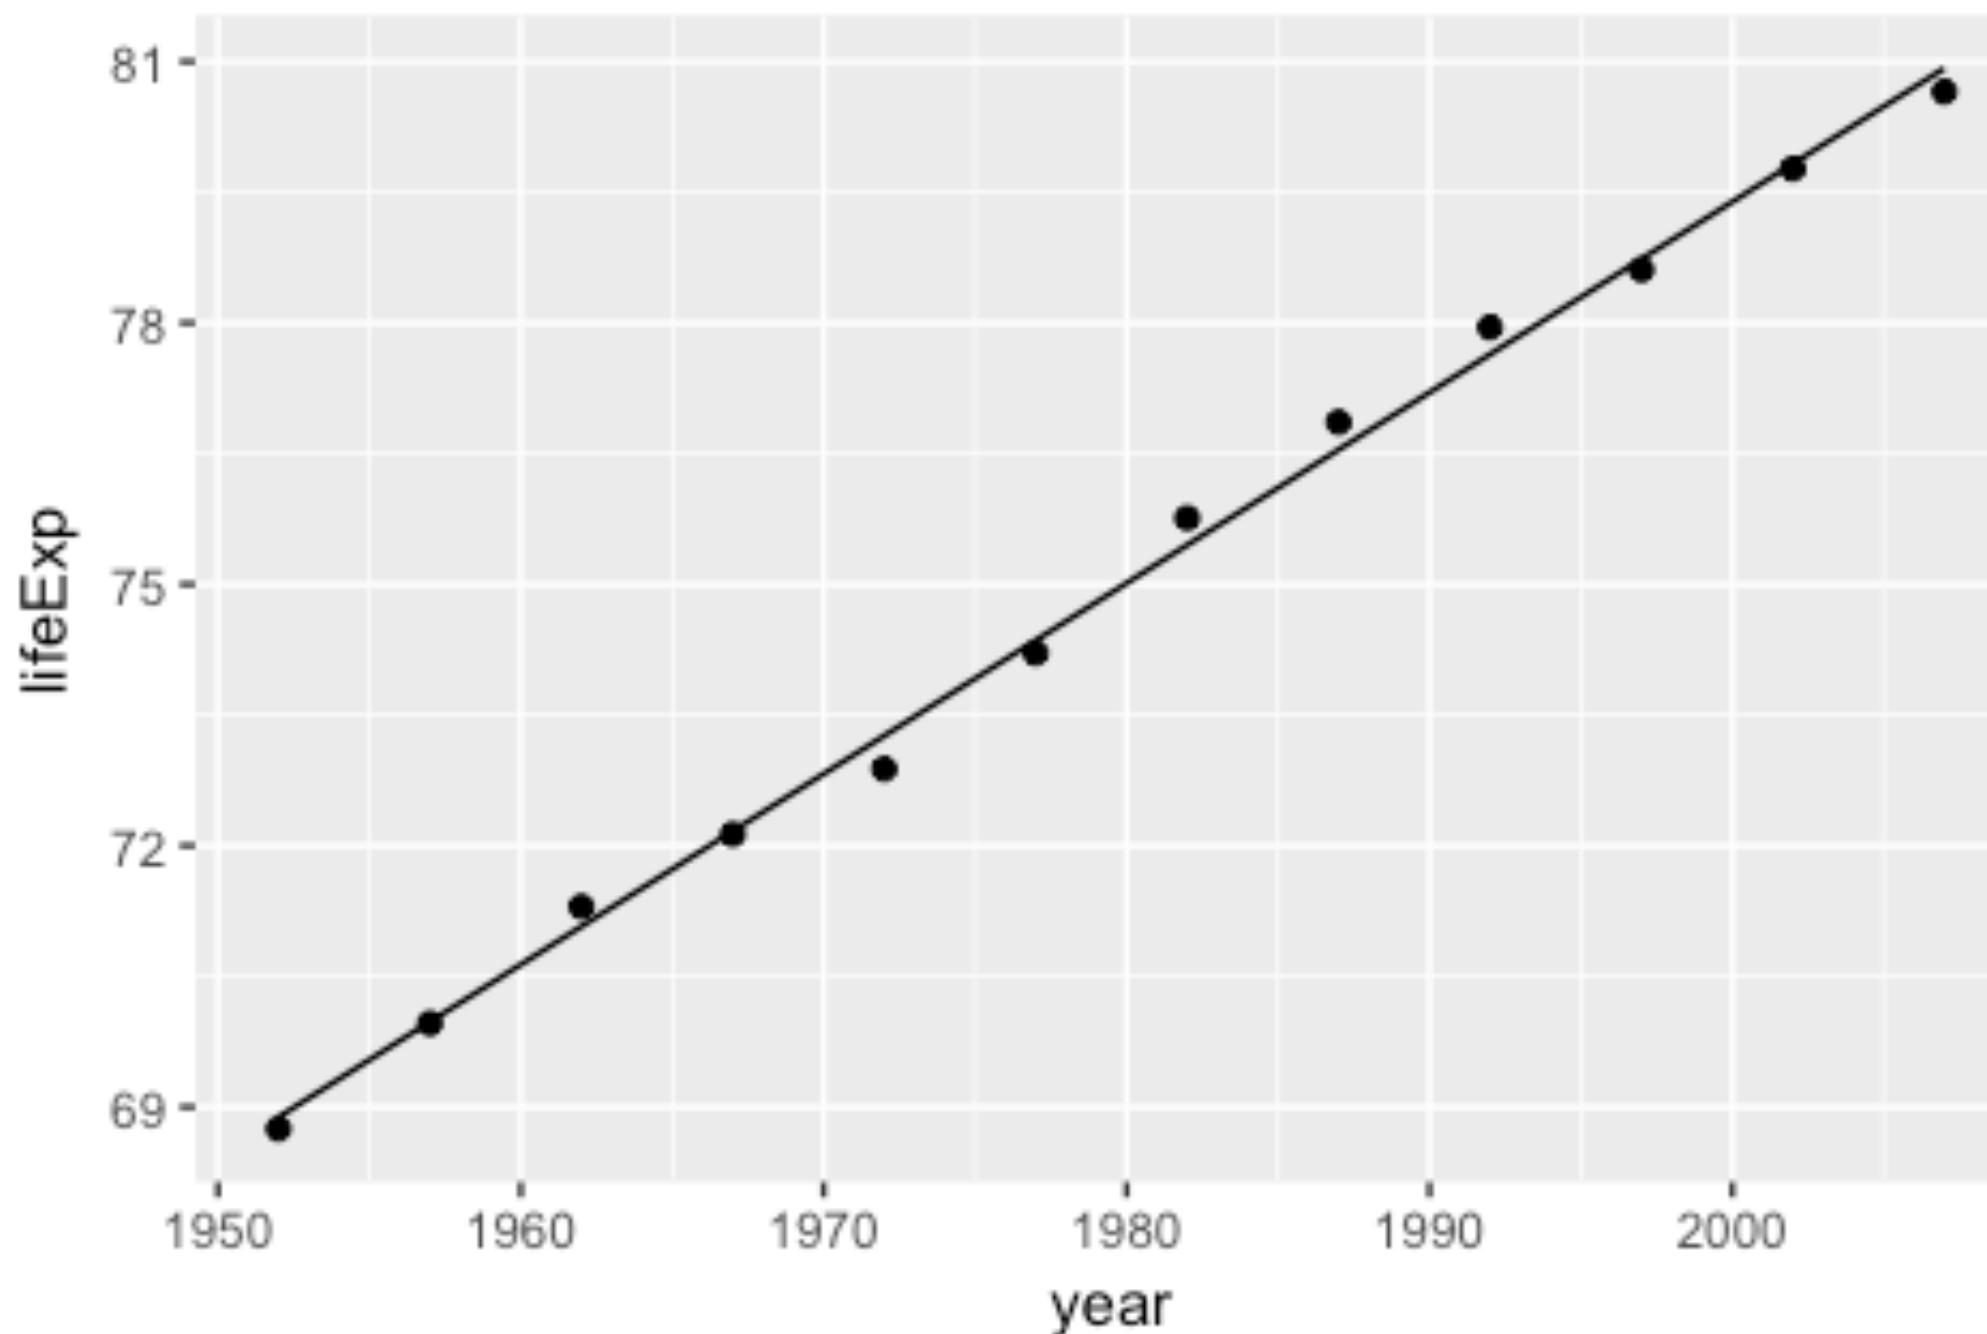
Fitted values

Residual



augment() useful for visualizing

```
canada_lm %>% augment() %>%  
  ggplot() +  
    geom_point(aes(x = year, y = lifeExp)) +  
    geom_line(aes(x = year, y = .fitted))
```



Recap



Use `glance()`, `tidy()`, and `augment()` to return model values in a data frame.

Your Turn 4

How has life expectancy changed in other countries?

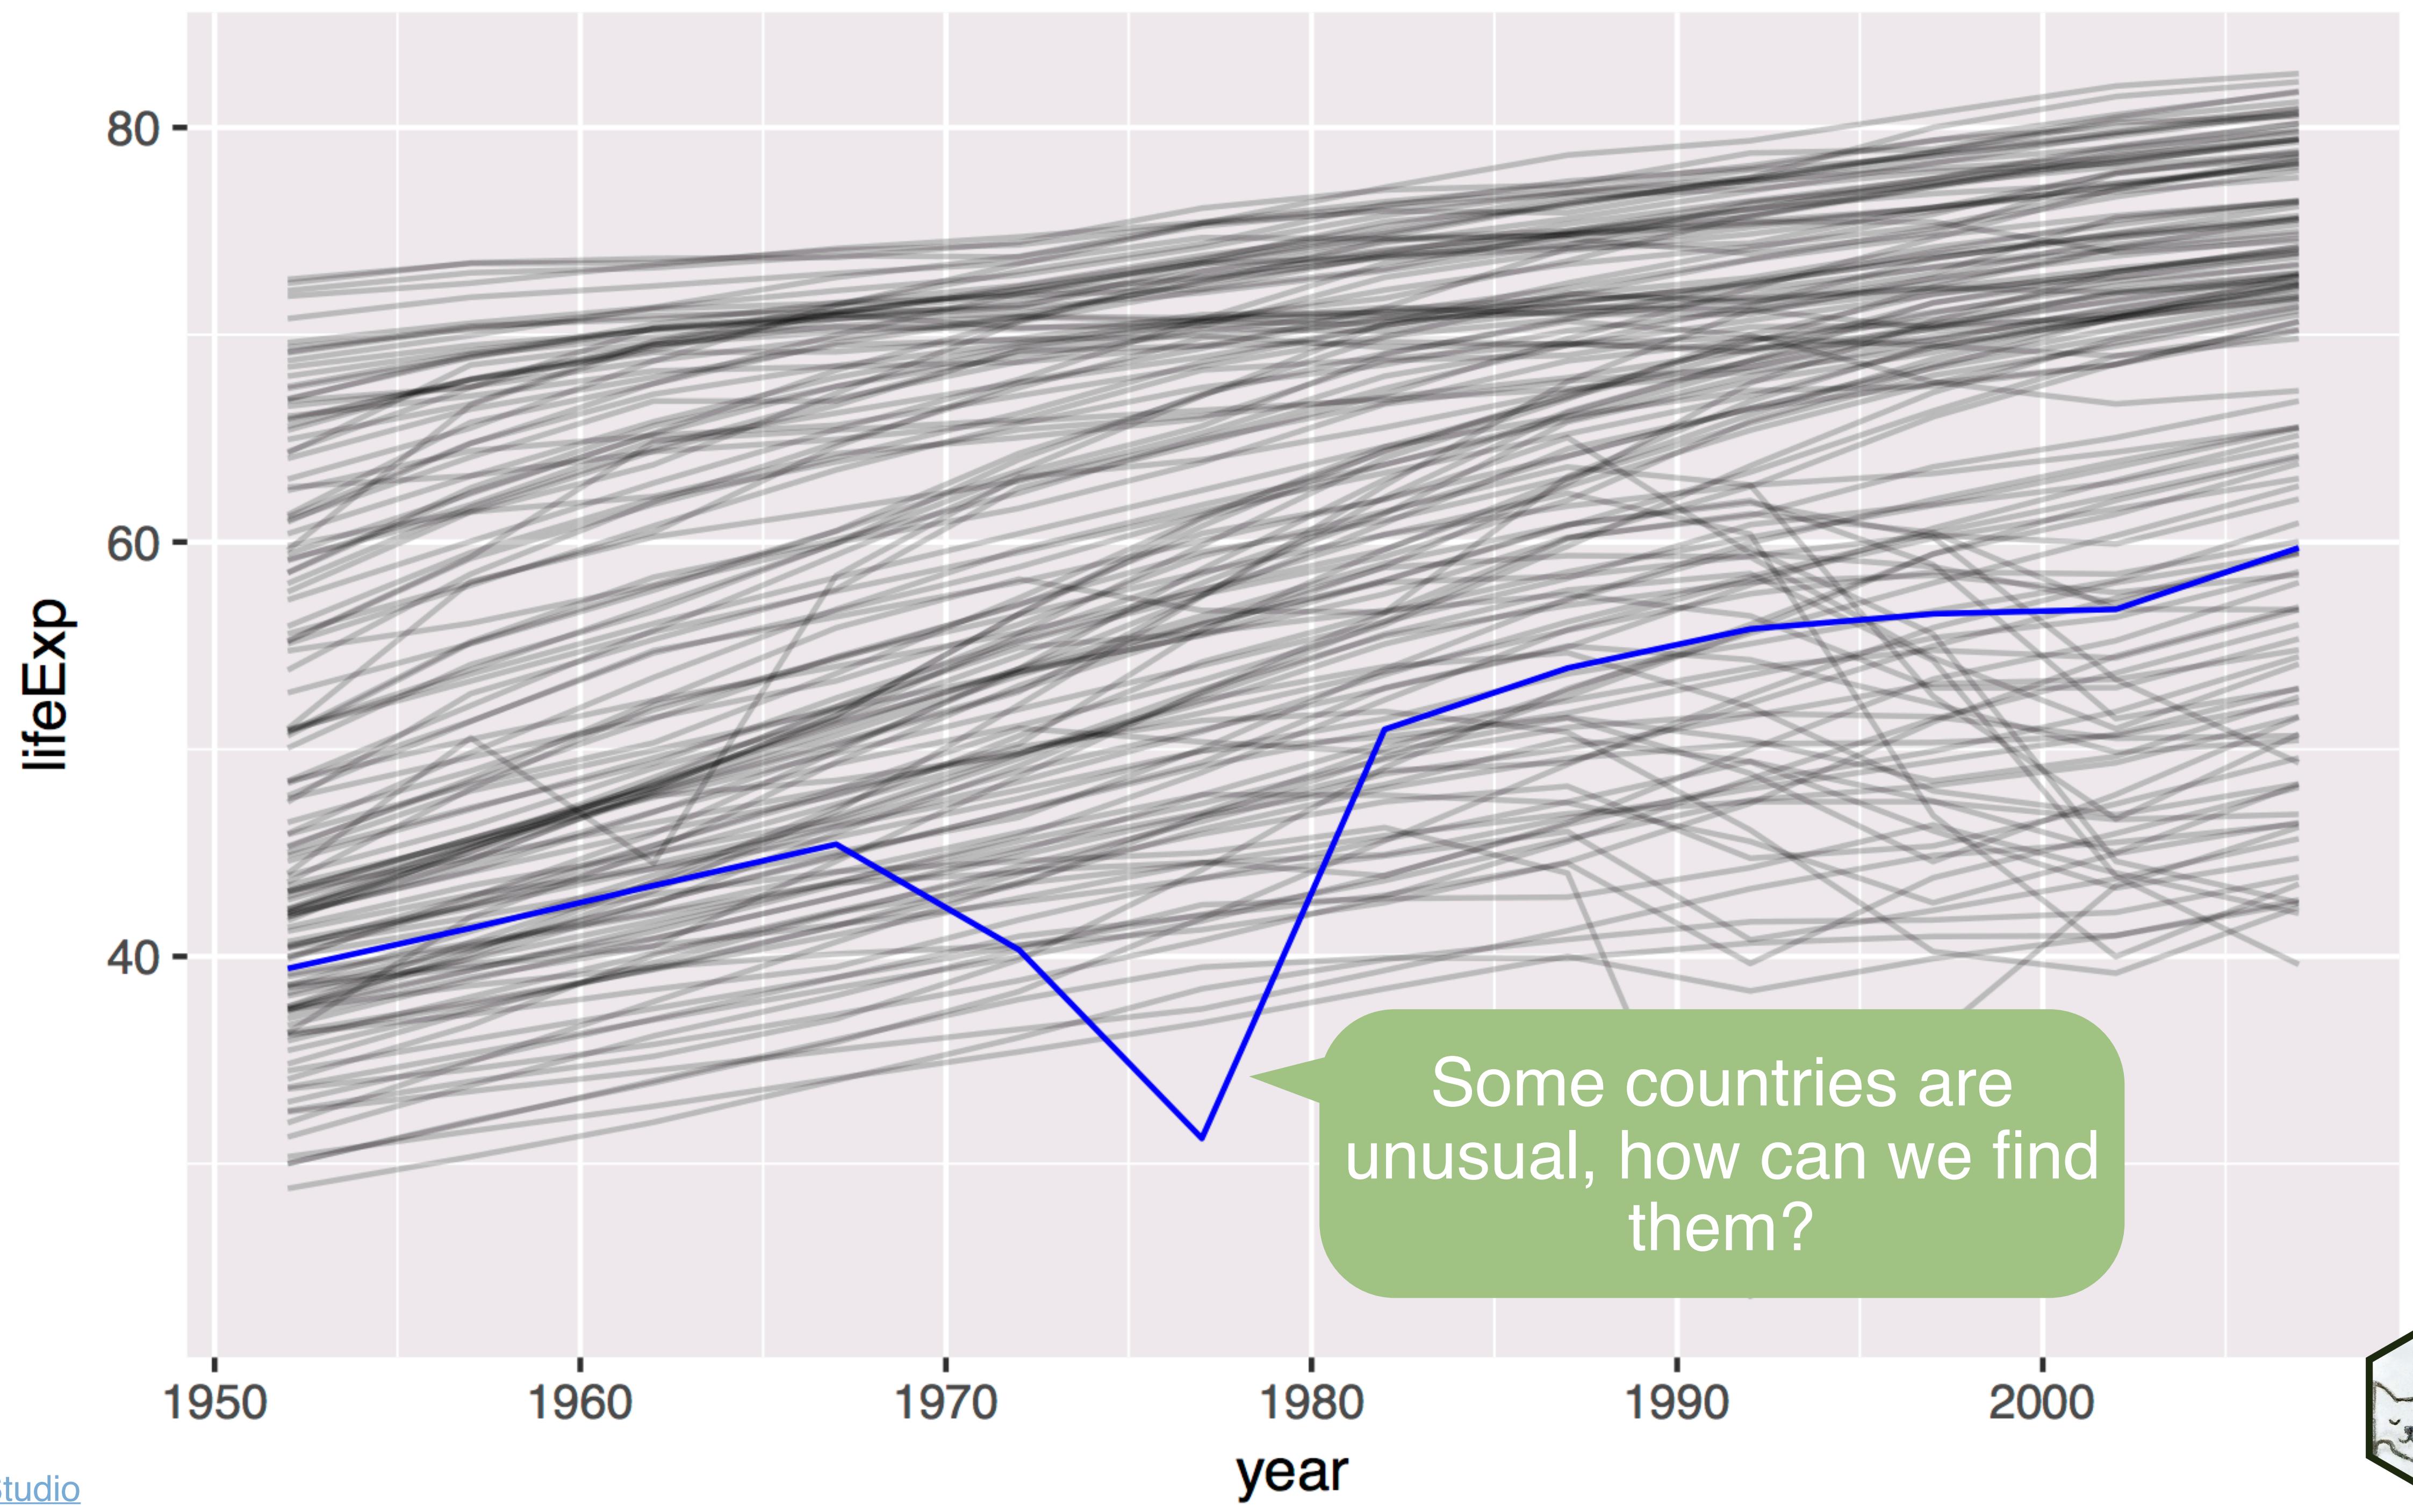
Make a line plot of lifeExp vs. year grouped by country. Set alpha to 0.2, to see the results better.



```
gapminder %>%  
  ggplot(mapping = aes(x = year, y = lifeExp, group = country)) +  
  geom_line(alpha = 0.2)
```



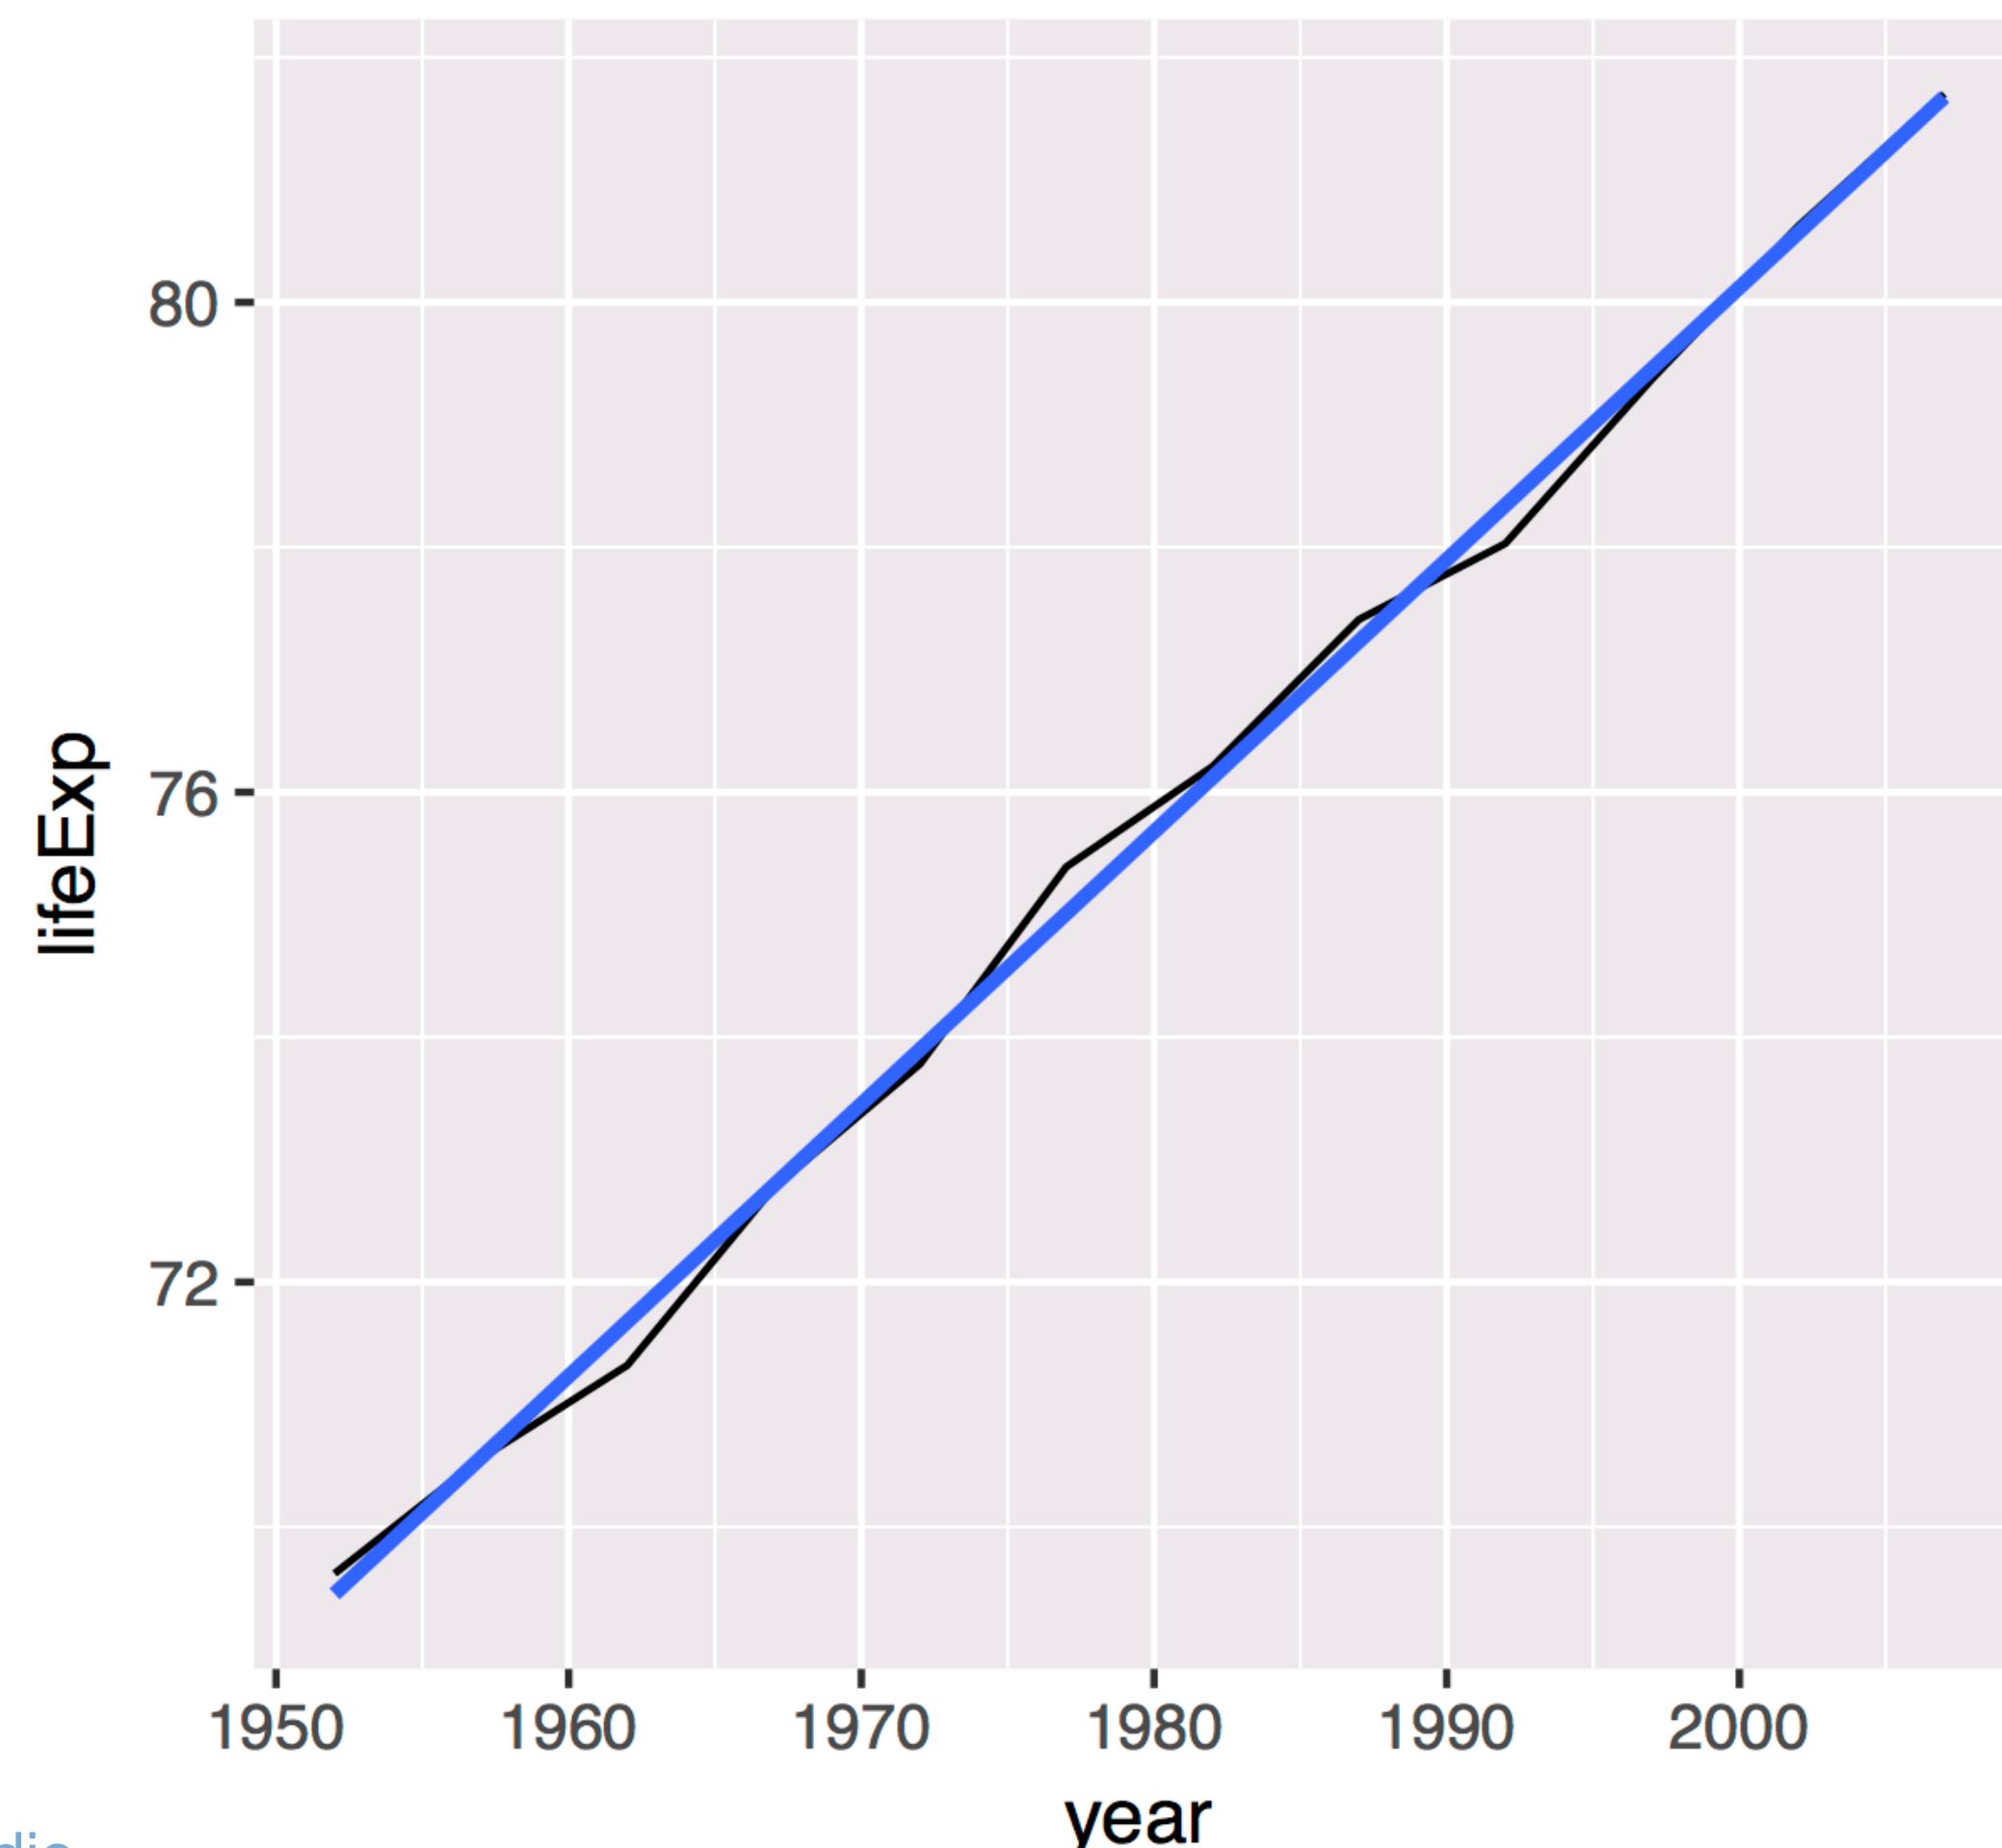




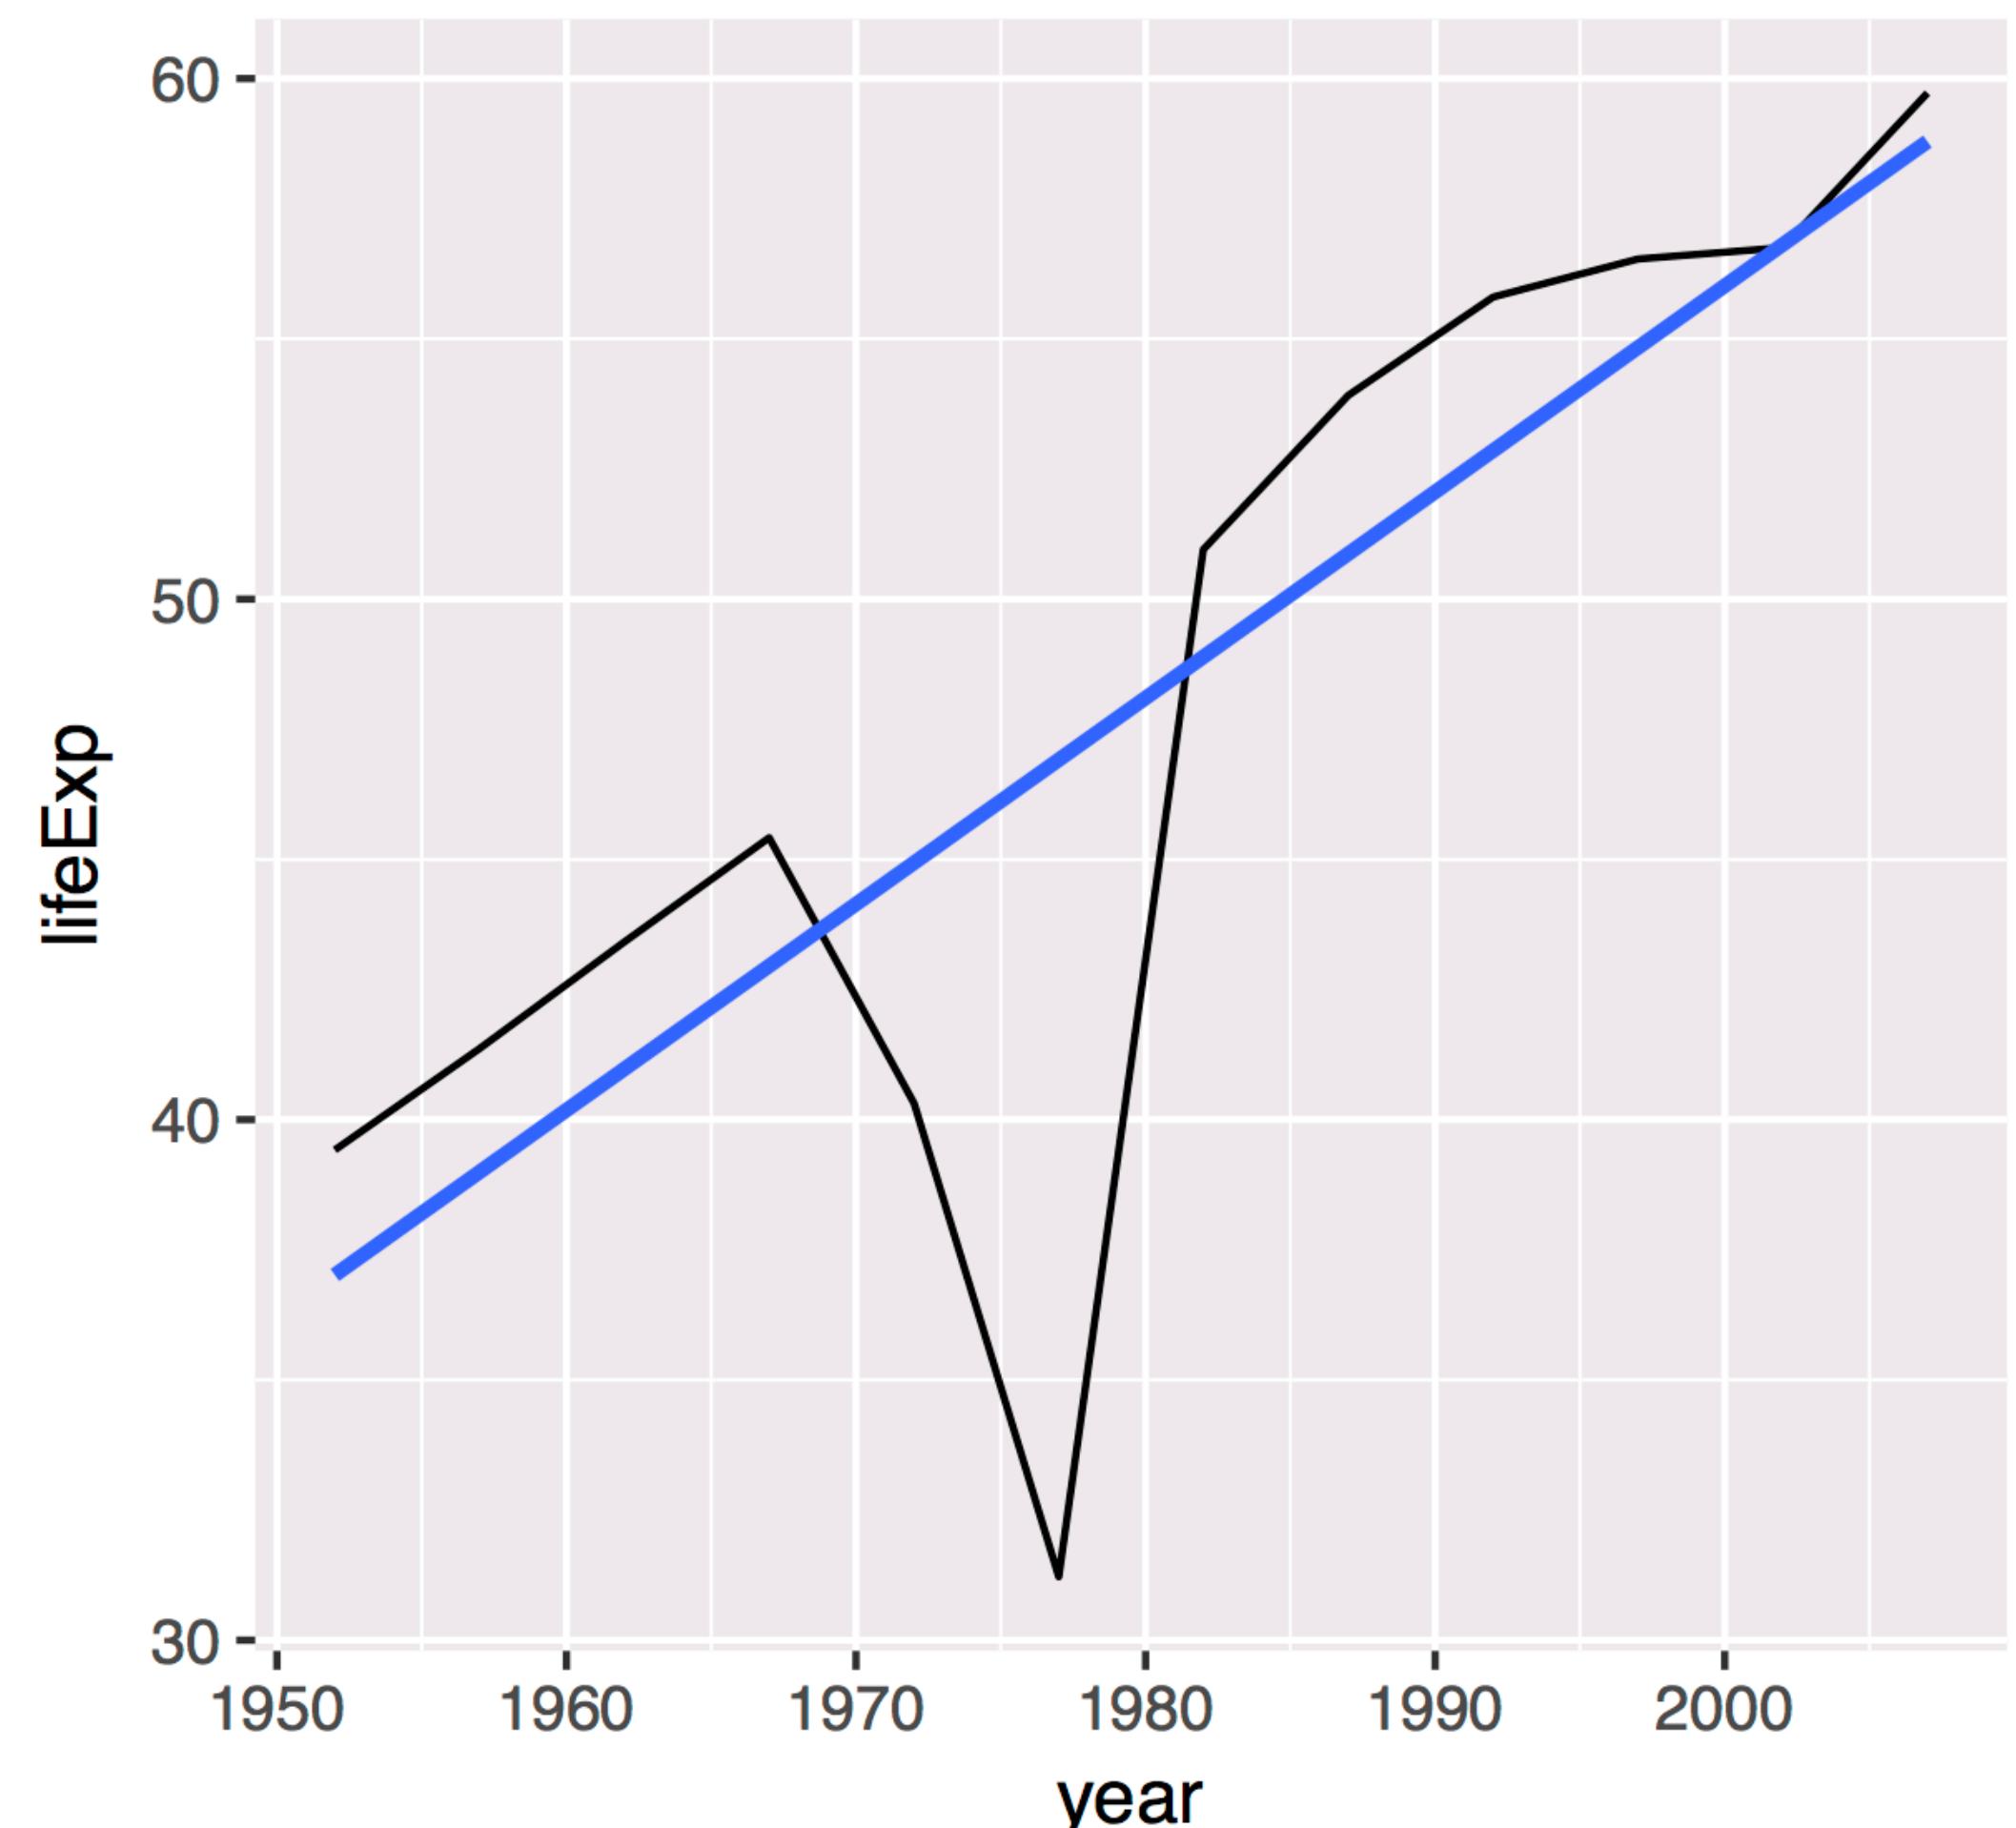
Idea 1

To quantify "linearity," fit a linear model, compare [r-squared](#).

Switzerland, R Squared = 0.99



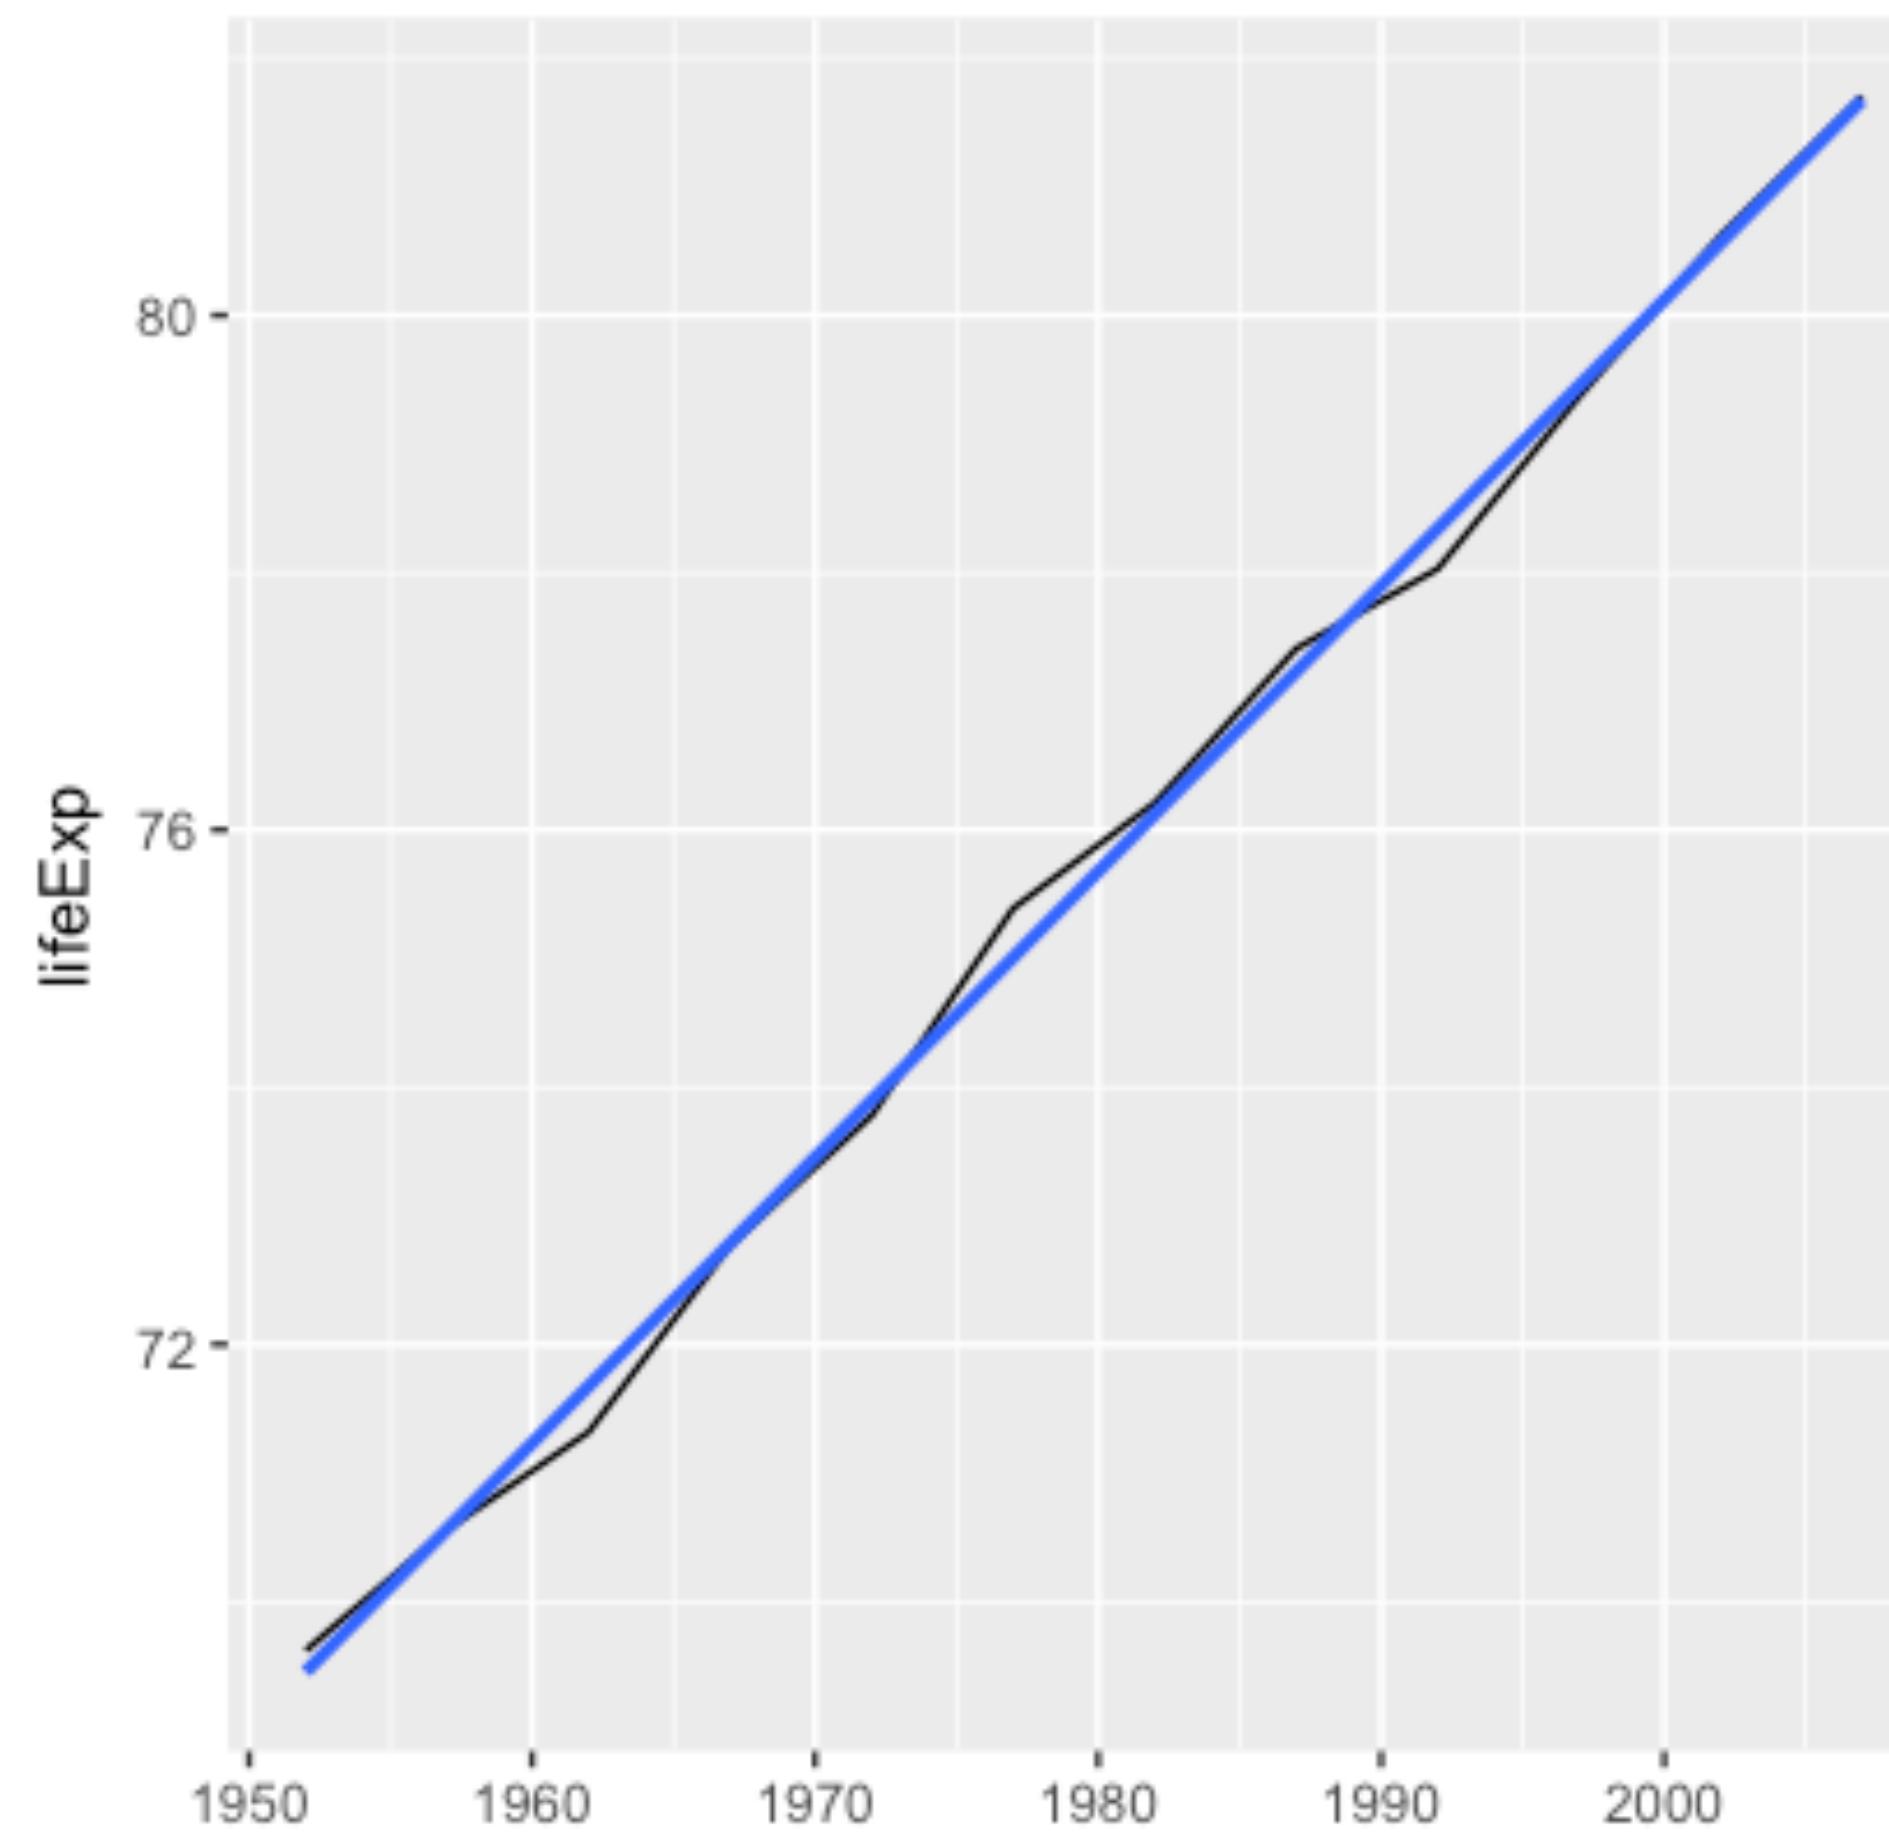
Cambodia, R Squared = 0.63



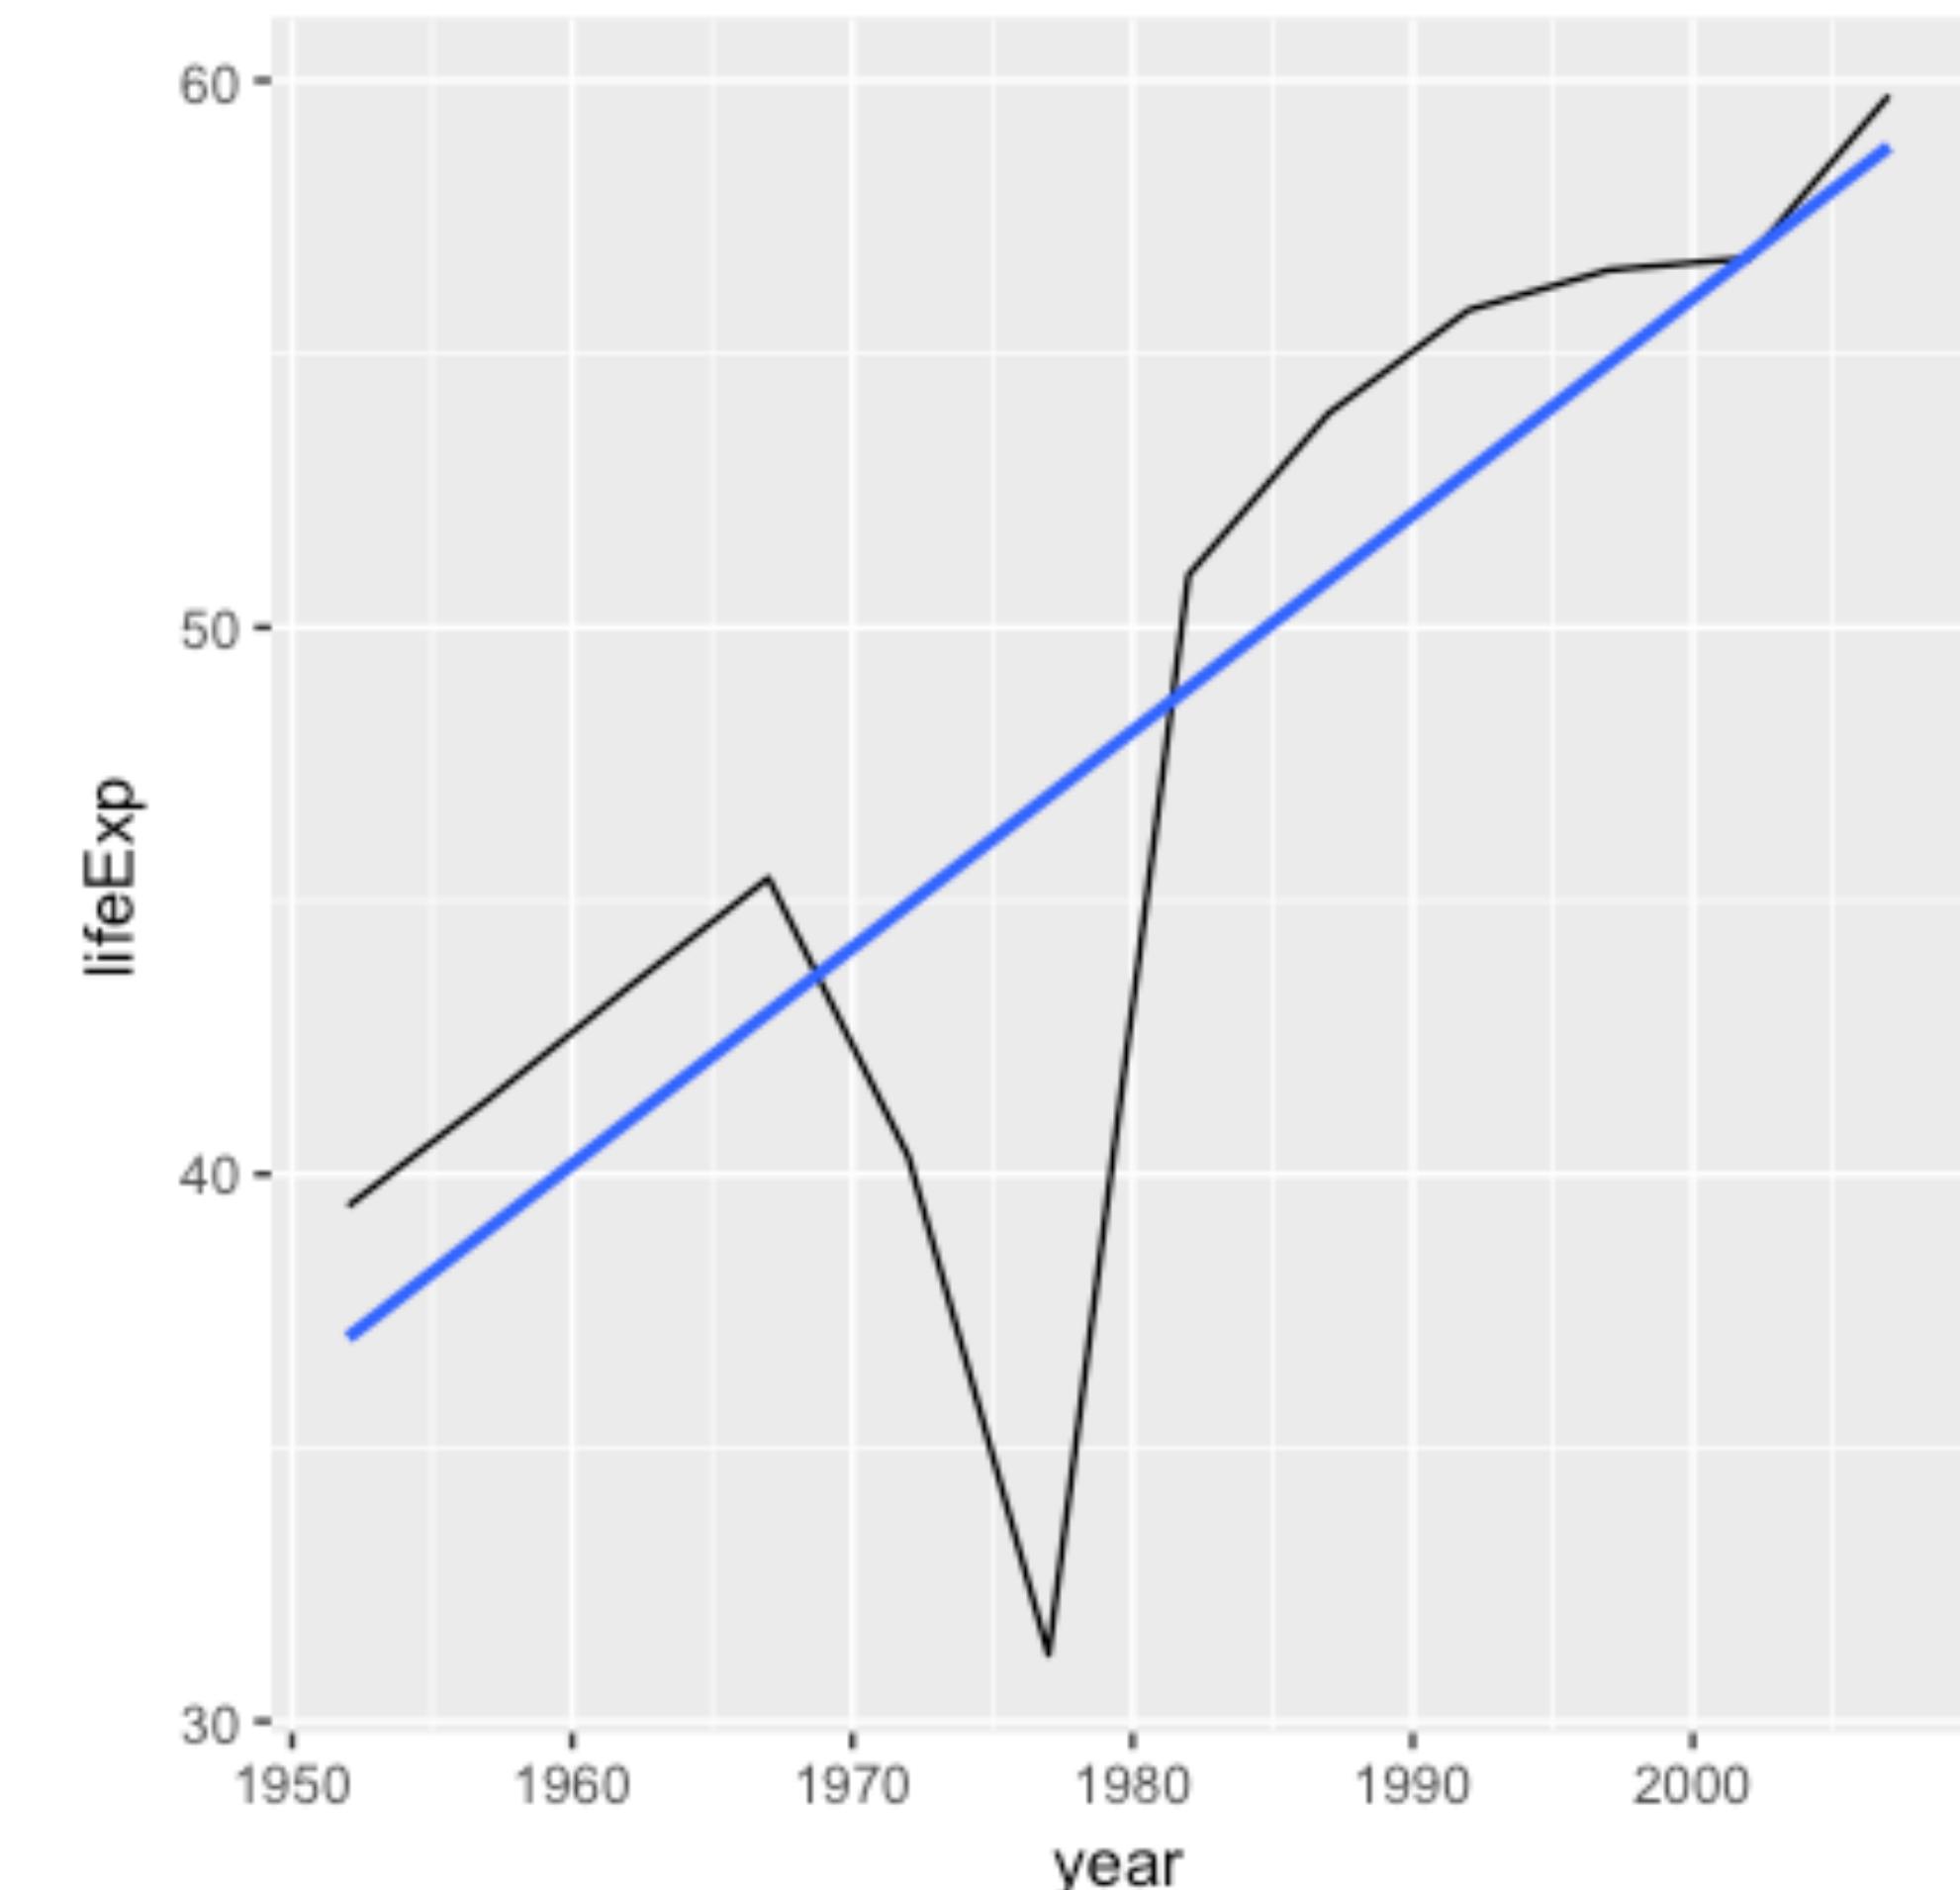
Idea 2

To quantify rate of change fit a linear model, extract **coefficient on year**.

Switzerland = 0.22 years/year



Cambodia = 0.40 years/year



Goal

Fit model, compute r.squared, collect coefficient **for every country.**

1. **dplyr** grouping toolkit
2. **purrr** toolkit and list columns